

Life cycle map

The NBN software development lifecycle as a cycle of eight stages, opened to task level. Each stage box opens into its tasks; each task says who acts, what to do, the colour of the delegation decision, the gate, and what would change the colour. The twelve D3 modules are pinned to the rows they govern.

Zac Clarkson (UX) · Source of truth: [docs/research/lifecycle-task-map.md](#) and [docs/reports/D3-sdlc-research.md](#) sections 5 and 6 · Departures from the D2 v2 baseline are proposals to Alessio

— How to read a box

COLOUR IS THE DELEGATION DECISION

Red A human decides, judges or approves.

Amber AI drafts or assists, a human approves.

Green AI executes, a human monitors.
Defensible only where a hook or an equivalent deterministic check stands behind it.

ACTOR IS WHO OR WHAT DOES THE TASK

 **Human.** A person decides or fires the command.

 **Agent.** The coding agent (Claude Code, Copilot agent) drafts, plans or executes.

 **Sub-agent.** A separate agent with its own context, used for review or a spike. It widens coverage; it is not an independent oracle.

 **Hook.** A deterministic check: git hook, CI job, design-system constraint, rollout controller. The only holder that makes green defensible.

GOVERNANCE TOUCHPOINTS

Token spend. Record generating time, reviewing time and tokens consumed for this task (module 6.9). The platform quota is the only hard stop and arrives without warning. On a row: **token spend**.

Code authorship. Authorship is four names: committer, agent co-author, non-requesting approver, accountable owner in the register (module 6.10). On a row: **code authorship**.

RULES THAT APPLY TO EVERY ROW

Anchor rule. Start from what is known to be real. Let AI generate only in the gaps deliberately left open. Research every generated hypothesis before treating it as real.

Exposure rule. Verification rigour scales with exposure: internal concept needs little, internal production tool needs a high standard, customer-facing production asset gets the full battery.

Review load. Diligence lands on the reviewer when the agent owns creation. Count review time beside generation time (modules 6.3, 6.7, 6.9). On a row: **review load**.

Green means a hook. Every green row names the deterministic check behind it. Where the check is not in force in our repository, the row says so.

Colour is never the only carrier of meaning on this map: a task row carries its colour as a band behind the id and the opened row states the word, following the accessibility rule the map itself sets in Stage 4.

— The cycle



Every stage opens red, generates amber or green, and closes on a red decision or gate. The dotted arrow from 8 back to 1 is 8.6 writing the learning back into the context 1.1 prompts from. Click a stage to open it. Stage colours were removed on 4 Sep because a stage banner adds nothing once the tasks carry colours; each box shows its profile and a tally of its task colours instead. Task 8.6 writes what production taught back into the context that Stage 1 prompts from, which is what makes this a cycle rather than a line.

The eight stages

1

Ideation and business case

Profile: leans intervention. 5 tasks.



WHERE THE BASELINE STOOD, AND WHAT THE TESTING FOUND

Baseline red with an amber first draft; sources silent. The framing and the decision are human; the divergent drafting between them is where the automation lives.

SHAPE OF THE STAGE

Brief becomes prompt, slip into automation for divergent options, return to a human decision.

ALSO IN THIS STAGE, HUMAN-LED WITH AI DRAFTING WHERE THE EVIDENCE RULE ALLOWS

Competitor and analogue scan, cost and benefit estimate, risk register first cut, success measures, sponsor identification, funding gate.

AS IS AT NBN TODAY

A product owner writes the case in a Word template, mostly alone, from the last similar project's numbers. AI, if used, is Copilot or a personal chat account with no record. Pain: the gap between first draft and a case the sponsor will read, and numbers that cannot be traced to a source. Hypothesis, verify in interviews.

Tasks, in order

1.1

Frame the brief



WHAT TO DO

1. Write the business context, the goals, what is troubling and what is on the horizon.
2. Get the sponsor to agree the brief describes the problem.
3. Treat the brief as the prompt: everything downstream runs on it.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

A written brief exists and the sponsor agrees it describes the problem.

WORKED EXAMPLE

Story A: the brief names missed appointments and the \$25 rebate as the pain to size.

WHO OR WHAT ACTS

Human Human

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. Framing is the human control surface for the whole stage.

1.2

Contextual inquiry



WHAT TO DO

1. Point the agent at public sources and at whatever internal material the account in use is cleared for (A4).
2. Ask for a scan that narrows from organisation to team to the specific process.
3. Strike out anything the agent asserted without a source before anyone reads it as fact.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Human has struck out anything the AI asserted without a source.

WORKED EXAMPLE

Base case: scan the technician fault-reporting process from public NBN and ACCC material, then confirm the team-level detail in interviews.

WHO OR WHAT ACTS

 **Agent** AI drafts the environment scan from public sources, human selects

COLOUR, AND WHAT WOULD MOVE IT

Amber Toward green only if a citation hook verifies every source before a human sees the scan. Toward red if the inquiry needs internal or personal data.

1.3 Generate divergent business-case options

 token spend

WHAT TO DO

1. Ask for three to five business cases with different framings, each with industry examples.
2. Choose between drafts instead of editing one into shape.
3. Ship nothing from this task.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Three to five options exist, none of them shipped anywhere.

WORKED EXAMPLE

The workshop drafted a full case (problem, solution, three benefits with metrics, effort, risk) in under a minute. Ask for that five times.

WHO OR WHAT ACTS

 **Agent** AI generates, human monitors

COLOUR, AND WHAT WOULD MOVE IT

Green Green because nothing ships (exposure rule). It becomes amber the moment one option is chosen and evidence is attached (1.4). Departure from the baseline, stated as a proposal.

GOVERNANCE TOUCHPOINTS

Token spend. Record generating time, reviewing time and tokens consumed for this task (module 6.9). The platform quota is the only hard stop and arrives without warning.

1.4 Back the preferred option with evidence



WHAT TO DO

1. Agent compiles evidence and cites every figure.
2. Run the citation-verification hook over the references (the workshop's model invented references).

WHO OR WHAT ACTS

 **Agent** AI compiles and cites, hook verifies citations, human runs interviews

3. Humans run the interviews and confirm volume figures with NBN before the case goes to a sponsor.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Every figure in the case has a verified source.

WORKED EXAMPLE

Story A: the rebate is public, so the agent sizes the case from it; the human confirms the missed-appointment volume with NBN.

COLOUR, AND WHAT WOULD MOVE IT

Amber The compilation moves toward green only when the citation hook is in force. The interviews stay red.

1.5

Refine, align stakeholders, sign off



WHAT TO DO

1. Refine the chosen case with the evidence attached.
2. Align stakeholders.
3. Record sponsor sign-off.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Sponsor sign-off recorded.

WORKED EXAMPLE

The agent's five options and the verified evidence are the input; the decision is the sponsor's.

WHO OR WHAT ACTS

Human Human

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. This is the decision the stage exists to make.

2 -

Requirements and discovery

Profile: mixed. 5 tasks.



WHERE THE BASELINE STOOD, AND WHAT THE TESTING FOUND

Baseline green for first-draft stories; slice 1 challenged it to amber. The task view dissolves the disagreement: drafting is green, the Definition of Ready gate is red.

SHAPE OF THE STAGE

Fix the anchors, draft in the gaps, verify against the Definition of Ready, commit to the repo.

ALSO IN THIS STAGE, HUMAN-LED WITH AI DRAFTING WHERE THE EVIDENCE RULE ALLOWS

Stakeholder interviews and contextual inquiry with real users, heuristic review of the current tool, non-functional requirements, data and privacy impact assessment, regulatory mapping (module 6.12), story mapping and release slicing, grooming cadence.

AS IS AT NBN TODAY

A business analyst writes stories in Jira from meeting notes and a Confluence page; out-of-scope lives in someone's head. Handoff is a document handed to the next role. Pain: scope creep from criteria never written down (the 35-file change) and re-explaining context. Hypothesis, verify in interviews.

Tasks, in order

2.1

Fix the anchor set



WHAT TO DO

1. Write down the known users, known tasks, known pain points and existing evidence.
2. Mark each as fact, not hypothesis.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

The anchor set is written down and marked as fact.

WORKED EXAMPLE

Base case: the technicians, the fault types, the current reporting steps and the known complaints.

WHO OR WHAT ACTS

Human Human

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. The anchor rule needs a human to say what is real.

2.2

Draft user stories from the anchor set

token spend

WHAT TO DO

1. Give the agent the anchor set and ask for stories that trace to it.

WHO OR WHAT ACTS

Agent AI generates, human monitors

2. Expect the drafts to be stylistically fine, less diverse than human ones, and light on acceptance criteria (slice 1, Stage 2).
3. Do not accept anything here; acceptance is 2.4.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Every story traces to an anchor or to a gap deliberately left open.

WORKED EXAMPLE

Story A: stories for the retailer raising a request, the dispatcher matching equipment, the technician confirming the window.

COLOUR, AND WHAT WOULD MOVE IT

Green Green because a draft story is internal and low exposure and the knowledge is not novel. The colour the literature argues about belongs to 2.4.

GOVERNANCE TOUCHPOINTS

Token spend. Record generating time, reviewing time and tokens consumed for this task (module 6.9). The platform quota is the only hard stop and arrives without warning.

2.3 Draft acceptance criteria and the out-of-scope list

 review load

WHAT TO DO

1. Ask for acceptance criteria and an explicit out-of-scope list per story.
2. Add what the agent did not think to ask for: accessibility and security criteria on every story.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Accessibility and security criteria present on every story.

WORKED EXAMPLE

A workshop team found the generated criteria had no accessibility requirements until prompted in (D2 v2 section 7).

WHO OR WHAT ACTS

 **Agent** AI drafts, human edits

COLOUR, AND WHAT WOULD MOVE IT

Amber Toward green if a template check refuses a story without accessibility and security criteria. Stays amber while the human is adding what the agent missed.

GOVERNANCE TOUCHPOINTS

Review load. Diligence lands on the reviewer when the agent owns creation. Count review time beside generation time (modules 6.3, 6.7, 6.9).

2.4 Verify against the Definition of Ready, prioritise, accept



WHAT TO DO

1. Walk each story through the Definition of Ready checklist.
2. Prioritise.
3. Product owner accepts into the backlog.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Story passes the DoR checklist; product owner accepts.

WHO OR WHAT ACTS

 **Human** Human

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. The DoR is a team agreement a person walks (Microsoft playbook).

WORKED EXAMPLE

Leon, 4 Sep: AI-forward for drafting, but at the end you verify against the Definition of Ready.

2.5

Commit accepted stories and criteria to the repo as Markdown



WHAT TO DO

1. Commit the accepted stories and criteria as a Markdown file on the branch.
2. Reference the file by name in every later prompt instead of pasting context.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

File is on the branch and referenced by name in later prompts.

WORKED EXAMPLE

Several workshop teams invented this independently (D2 v2 section 7.3). It is the context lever for this stage.

WHO OR WHAT ACTS

Hook Hook (CI check that the file exists and parses)

COLOUR, AND WHAT WOULD MOVE IT

Green Green while the CI check exists. Without it the task is a convention.

MODULES PINNED HERE

6.7 Collaboration with an agent in the loop

3

Solution design and architecture

Profile: leans intervention, with the largest amber-if-context lever in the lifecycle. 5 tasks.



WHERE THE BASELINE STOOD, AND WHAT THE TESTING FOUND

Baseline red as delivered, amber if context improves; both halves confirmed. The lever is task 3.1: the better the documented context, the more of 3.2 and 3.5 the AI can carry.

SHAPE OF THE STAGE

Document the as-is, generate options against the approved list, decide as an ADR, threat model, spike the riskiest unknown.

ALSO IN THIS STAGE, HUMAN-LED WITH AI DRAFTING WHERE THE EVIDENCE RULE ALLOWS

Data model and schema, integration and API contract design (story B's contract starts here), non-functional design, build-versus-buy against the approved list, cost estimate, design review sign-off, fitness functions for CI.

AS IS AT NBN TODAY

An architect draws the design against a partly out-of-date infrastructure map and records the decision on a Confluence page reviewed in a meeting. The estate is decades old in many languages. Pain: rediscovering the as-is instead of designing the to-be. Hypothesis, verify in interviews.

Tasks, in order

3.1

Document the existing system



WHAT TO DO

1. Point the agent at the repos, the infrastructure map and the existing patterns.
2. Have it draft the context document: infrastructure, APIs, patterns.
3. Correct it, then keep it in the repo where the next generation reads from it.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

A context document exists in the repo that a new engineer could onboard from.

WORKED EXAMPLE

The brand-colours fix was to have the model learn the design system and write it into the instructions file. 3.1 is that fix applied to architecture.

WHO OR WHAT ACTS

 **Agent AI** reads and drafts, human corrects

COLOUR, AND WHAT WOULD MOVE IT

Amber This task is what moves the rest of the stage. Better context makes 3.2 and 3.5 more reliable. It is the precondition for any amber-to-green move (slice 1 recommendation 4).

3.2

Generate candidate architectures against the approved tool list



WHAT TO DO

1. Load the context document and the approved tool list.
2. Ask for three to five options, each naming its trade-offs.
3. Select; do not accept the first.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Every option uses only approved tools and names its trade-offs.

WORKED EXAMPLE

Leon's example: a page with a map view needing credentials; the approved list names one mapping library, so the suggestion uses it.

WHO OR WHAT ACTS

Agent AI drafts three to five options, human selects

COLOUR, AND WHAT WOULD MOVE IT

Amber Only as good as the context it was pointed at (D1 section 3). Stays amber; a check that rejects unapproved tools would harden the gate; the choice stays with a person.

3.3

Trade study and architecture decision record



WHAT TO DO

1. Run the trade study on the options.
2. Write the ADR and merge it.
3. Update the decision log.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

ADR merged; decision log updated.

WORKED EXAMPLE

This repo already has one: docs/adr/001-stack.md.

WHO OR WHAT ACTS

Human Human

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it.

3.4

Threat model, security and privacy requirements



WHAT TO DO

1. Walk the threat model with the AI checklist as an assist.
2. Have someone other than the author review it.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Threat model reviewed by someone other than the author.

WHO OR WHAT ACTS

Human Human, AI-assisted checklist

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. Microsoft's SDL puts it in design with independent review.

WORKED EXAMPLE

Story A: the appointment data path and the retailer portal boundary are in scope; no production data in a prompt on an account that is not cleared for it (A4).

3.5

Technical spike on the riskiest unknown

 token spend

WHAT TO DO

1. Scope one question and a time-box.
2. Let the agent go and find out.
3. Log three numbers: generating time, reviewing time, tokens consumed (module 6.9).
4. Read the spike log and decide.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Spike log answers the question it was scoped for; time-box respected.

WORKED EXAMPLE

The team's own INF-6a pattern.

WHO OR WHAT ACTS

 **Sub-agent** AI agent runs the spike, human reads the log

COLOUR, AND WHAT WOULD MOVE IT

Amber Stays amber. The token register lives here; the platform quota is the only hard stop.

GOVERNANCE TOUCHPOINTS

Token spend. Record generating time, reviewing time and tokens consumed for this task (module 6.9). The platform quota is the only hard stop and arrives without warning.


MODULES PINNED HERE

6.9 Token accounting and cost

4

UX design and prototyping

Profile: mixed, and deliberately split. 8 tasks.



WHERE THE BASELINE STOOD, AND WHAT THE TESTING FOUND

Baseline green journey mapping, red brand and design; the green half challenged to amber, the red half sharpened by both supervisor reviews into a deterministic rule. The team's own stage.

SHAPE OF THE STAGE

Anchor, hypothesise in the gaps, skeleton fast, enforce the design system with a hook, critique with a human, test with real users.

ALSO IN THIS STAGE, HUMAN-LED WITH AI DRAFTING WHERE THE EVIDENCE RULE ALLOWS

Content and microcopy, information architecture, interaction states and error handling, motion policy, localisation and plain-English review, token maintenance, design QA against the build, screenshot-as-specification handoff.

AS IS AT NBN TODAY

A designer works in a design tool against the brand system and hands screens to engineering at the end. Prose is a poor input for spatial decisions; generated designs drift off the system as context fills; the developer sits idle meanwhile. NBN runs Siteimprove and treats WCAG 2.2 AAA as part of done. Best-evidenced as-is (Group 1 designers spoke).

What we chose not to do, as design. Because accessibility is a hard requirement at NBN, the design system the agent generates against has less motion, no auto-dismissing pop-ups, no information carried by colour alone, and a keyboard path through every flow. Those choices are enforced in 4.5, not reviewed in 4.6. This map follows its own rule: every colour also carries its word.

Tasks, in order

4.1

Fix the anchor set



WHAT TO DO

1. Write down known users, historical evidence, known tasks and pains as fact.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Anchors written down as fact.

WORKED EXAMPLE

Story A: dispatchers, technicians, the known missed-appointment causes.

WHO OR WHAT ACTS

Human Human

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it.

4.2

Hypothesis journey map and proto-personas in the gaps



WHAT TO DO

1. Let the agent draft only in the gaps the anchor set leaves open.
2. Attach a research task to every generated hypothesis.
3. Nothing generated is real until researched.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Every generated hypothesis has a research task attached.

WORKED EXAMPLE

NN/g hypothesis maps and Leon's anchor rule are the same instruction from two directions.

WHO OR WHAT ACTS

 **Agent AI** drafts, human researcher validates

COLOUR, AND WHAT WOULD MOVE IT

Amber Toward red if hypotheses start being treated as findings. Never green. A hypothesis map is a research plan.

4.3

Skeleton screens straight from the acceptance criteria



token spend

WHAT TO DO

1. Generate rough screens for every story in the sprint from the acceptance criteria file (2.5).
2. Keep them rough. They are things to react to.
3. Move on to the next page while the engineer works the first.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Rough screens exist for every story in the sprint, none of them polished.

WORKED EXAMPLE

The Tuesday team kept its engineer busy on page one while the design pair moved to page two (D2 v2 section 7.1).

WHO OR WHAT ACTS

 **Agent AI** generates, human monitors

COLOUR, AND WHAT WOULD MOVE IT

Green Green because it is exploration. It turns amber the moment a skeleton is treated as a deliverable.

GOVERNANCE TOUCHPOINTS

Token spend. Record generating time, reviewing time and tokens consumed for this task (module 6.9). The platform quota is the only hard stop and arrives without warning.

MODULES PINNED HERE

6.7 Collaboration with an agent in the loop

4.4

Brand and design-system decisions



WHAT TO DO

1. Decide the brand and system rules as humans.
2. Load the design tokens into context as a constraint, not a description, before any generation.

WHO OR WHAT ACTS

 **Human** Human; the design system is the rule set

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Design tokens loaded into context before any generation.

WORKED EXAMPLE

The brand-colours incident: the colours came from training data, and nobody could say where from until asked.

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. Generation is the wrong tool for enforcing a deterministic system (31 Aug). Edge case: a new device type gets green exploration first, then humans codify the findings back here.

MODULES PINNED HERE

[6.5 Design-system enforcement in build](#)

4.5

Brand compliance check by overlay diff



WHAT TO DO

1. Generate the layout.
2. Load the real design-system CSS after it so its rules override.
3. Whatever visibly shifts (fonts, colours, spacing) is the violation. Run as a Playwright visual comparison plus an axe scan.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

No visible shift, or every shift is explained.

WORKED EXAMPLE

The one green in the lifecycle that sits next to a red decision without contradiction.

WHO OR WHAT ACTS

Hook Hook (design-system CSS overrides generated CSS; visible shift fails)

COLOUR, AND WHAT WOULD MOVE IT

Green Green only while the check runs in CI. Playwright is not yet in this repo (recorded gap), so today the check exists on paper only.

MODULES PINNED HERE

[6.5 Design-system enforcement in build](#)

4.6

Usability and accessibility critique

[review load](#)

WHAT TO DO

1. Run the AI accessibility pass and the axe scan as triage.
2. A human expert dispositions every finding.
3. Encode a repeated finding as a skill so later pages are generated against it.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Human has dispositioned every AI finding; WCAG target met.

WORKED EXAMPLE

Team 6 turned page one's findings into an accessibility skill (D2 v2 section 7.3).

WHO OR WHAT ACTS

Agent AI first-pass triage feeding a human expert

COLOUR, AND WHAT WOULD MOVE IT

Amber AI covers roughly two thirds of expert-found errors with meaningful false positives (slice 1). Moves toward green only for the findings a skill has turned into a generation rule; the expert stays.

GOVERNANCE TOUCHPOINTS

Review load. Diligence lands on the reviewer when the agent owns creation. Count review time beside generation time (modules 6.3, 6.7, 6.9).

MODULES PINNED HERE

6.5 Design-system enforcement in build

6.7 Collaboration with an agent in the loop

4.7

Inclusion decisions, including First Nations inclusive-design policy



WHAT TO DO

1. Codify the policy in the system prompt and tool chain before the stage starts.
2. Treat generated artifacts for First Nations audiences as hypotheses until validated with the community.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Policy codified in the system prompt and tool chain before design starts.

WORKED EXAMPLE

Regulatory rule placed ahead of design (module 6.12).

WHO OR WHAT ACTS

Human Human; codified as a hard limit before design starts

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. Stretch RAP VI 2026 to 2029 and the First Nations Digital Inclusion Plan bind the design, and the AI too.

MODULES PINNED HERE

6.5 Design-system enforcement in build

6.12 Regulatory codification for NBN Co

4.8

Test with real users



WHAT TO DO

1. Test with real users.
2. Record the findings.
3. Do not substitute synthetic users.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Findings from real users recorded; synthetic users not used as a substitute.

WORKED EXAMPLE

Design Sprint Friday is real customers for the same reason.

WHO OR WHAT ACTS

Human Human

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. Synthetic users idealise and fabricate (slice 1, Stage 4).

5

Development planning

Profile: mixed. 5 tasks.



WHERE THE BASELINE STOOD, AND WHAT THE TESTING FOUND

Baseline amber, confirmed. The most important row is the one the workshop weighted most: writing the issue properly, because the out-of-scope list prevents the thirty-five-file change.

SHAPE OF THE STAGE

AI proposes the slicing and priorities, a human writes the issue and commits the plan, a hook pushes the issues.

ALSO IN THIS STAGE, HUMAN-LED WITH AI DRAFTING WHERE THE EVIDENCE RULE ALLOWS

Capacity check, risk and dependency review, spike identification, Definition of Done, environment and access readiness (tenant and account boundaries cost the workshop time), sprint goal, communication plan for affected retailers (story B).

AS IS AT NBN TODAY

A delivery lead runs planning in Jira; stories are estimated in a session, assigned by who is free, dependencies found when someone is blocked. Out-of-scope is rarely written. Hypothesis, verify in interviews.

Tasks, in order

5.1

Slice the work into minimal, single-purpose issues



WHAT TO DO

1. Ask the agent to slice the accepted stories into single-purpose issues.
2. Reject any issue that bundles more than one purpose.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

No issue bundles more than one purpose.

WORKED EXAMPLE

Story A: appointment-window logic and technician matching are separate issues, never one.

WHO OR WHAT ACTS

 **Agent** AI proposes the slicing, human adjusts

COLOUR, AND WHAT WOULD MOVE IT

Amber Stays amber. Agentic PRs bundle purposes at over three times the human rate (slice 1, Stage 5), so the small-batch rule needs a person.

5.2

Prioritise, assign to roles, flag dependencies and complexity



WHAT TO DO

1. Have the agent draft MoSCoW priority, role assignment, dependencies and complexity.
2. Edit against capacity and what the team knows.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Every issue has a priority, an owner role and its dependencies listed.

WORKED EXAMPLE

The workshop demonstrated this end to end into a two-week plan (D2 v2 section 4.1).

WHO OR WHAT ACTS

 **Agent** AI drafts, human edits

COLOUR, AND WHAT WOULD MOVE IT

Amber Stays amber; the plan is only as good as the meeting's knowledge.

5.3 Write each issue properly



WHAT TO DO

1. Write the acceptance criteria and the out-of-scope list yourself.
2. Make the criteria testable.
3. Point the issue at the stories file from 2.5.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Out-of-scope list present; criteria testable.

WORKED EXAMPLE

The out-of-scope list is what prevents the 35-file change (D2 v2 section 4.2).

WHO OR WHAT ACTS

 **Human** Human (the instructing role)

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. This is the human's control surface over everything that follows: a good issue gives a good outcome.

MODULES PINNED HERE

[6.1 Issue to branch](#)

[6.11 Prompt practice: written, stored, versioned, shared](#)

5.4 Commit the sprint plan



WHAT TO DO

1. Team and product owner agree the plan.
2. Check capacity.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Team and product owner agree the plan; capacity checked.

WHO OR WHAT ACTS

 **Human** Human

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it.

5.5 Push the issues into the repo host



WHAT TO DO

1. Push the issues through the template.

WHO OR WHAT ACTS

 **Hook** Hook (issue template validation)

2. Let the template check refuse an issue missing criteria or an out-of-scope list.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Issues exist, templated, linked to the stories file from 2.5.

COLOUR, AND WHAT WOULD MOVE IT

Green

Green while the template validation exists. Mechanical.

6

Development and build

Profile: leans autonomous, with a red gate at the end. 5 tasks.



WHERE THE BASELINE STOOD, AND WHAT THE TESTING FOUND

Baseline amber, confirmed as the best-evidenced stage. Slice 2 modules 1 to 3 carry the operational detail.

SHAPE OF THE STAGE

Agent plans, agent builds behind hooks, agent opens the PR with provenance, a human who did not write it approves.

ALSO IN THIS STAGE, HUMAN-LED WITH AI DRAFTING WHERE THE EVIDENCE RULE ALLOWS

Local setup and environment parity, coding standards, secret handling and the no-NBN-data rule, dependency selection against the approved list, migrations, feature flag wiring, docs in the same PR, pairing where the change is rebate-relevant.

AS IS AT NBN TODAY

A developer picks up a Jira ticket, asks the analyst what it meant, branches, codes, writes some tests, opens a PR that waits. Not every developer is on GitHub. AI, if used, is Copilot or a personal chat with no record. Pain: the wait for review, and the decision load on the reviewer once AI owns creation. Hypothesis, verify in interviews.

Tasks, in order

6.1

Agent reads the issue, confirms each acceptance criterion, plans



WHAT TO DO

1. Point the agent at the issue and the acceptance-criteria file.
2. Have it restate each criterion and produce a plan.
3. Confirm or redirect before any code is written; this is the last cheap point to redirect.

WHO OR WHAT ACTS

 **Agent** AI plans, human confirms the plan

COLOUR, AND WHAT WOULD MOVE IT

Amber

Stays amber. The plan step escalates to a stronger model and is where redirection is cheapest.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Plan accepted before any code is written.

WORKED EXAMPLE

Base case: the fault-reporting issue, planned against its criteria file.

MODULES PINNED HERE

6.1 Issue to branch

6.2

Branch, implement, commit



token spend code authorship

WHAT TO DO

1. Agent branches from main, implements, commits.
2. Pre-commit hooks run lint and format; commit-msg enforces Conventional Commits.
3. One logical change per commit is a human rule the reviewer checks.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Every commit passes the pre-commit hooks.

WORKED EXAMPLE

Slice 2 module 2: how a commit is done.

WHO OR WHAT ACTS



Agent AI executes; hooks enforce lint, format, conventional commit, secret scan

COLOUR, AND WHAT WOULD MOVE IT

Green

Green only because hooks stand behind it. In this repo the prettier step ends in `|| true`, so format is advisory (recorded gap).

GOVERNANCE TOUCHPOINTS

Token spend. Record generating time, reviewing time and tokens consumed for this task (module 6.9). The platform quota is the only hard stop and arrives without warning.

Code authorship. Authorship is four names: committer, agent co-author, non-requesting approver, accountable owner in the register (module 6.10).

MODULES PINNED HERE

6.2 How a commit is done

6.3

Open a draft PR with description and provenance



code authorship

WHAT TO DO

1. Agent opens a draft PR through the template.
2. The PR carries the co-author trailer and the model and prompt record (the confidence card: plan, assumptions, alternatives, edge cases).

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

PR carries Co-Authored-By and the model and prompt record.

WHO OR WHAT ACTS



Agent AI executes; hook checks the provenance trailer

COLOUR, AND WHAT WOULD MOVE IT

Green

Amber in practice today: no hook checks the trailer and no commit on our main carries one (slice 3 module 4). A commit-msg check that refuses an agent commit without a trailer would make it green in an afternoon.

WORKED EXAMPLE

Story A: a missed-window rebate dispute a year later will ask what changed and who approved it (A3).

GOVERNANCE TOUCHPOINTS

Code authorship. Authorship is four names: committer, agent co-author, non-requesting approver, accountable owner in the register (module 6.10).

MODULES PINNED HERE

6.3 How branches become merges

6.10 Liability, provenance and code authorship

6.4

Human code review and merge approval

 code authorship review load

WHAT TO DO

1. Review in this order: does it match how we already do this, is it one change, did it update what it touched, and only then correctness.
2. Approve only if you did not request or author the change.
3. Merge on CI green.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Approved by someone other than the author; CI green.

WORKED EXAMPLE

The last human decision before production in this repository (module 6.6).

WHO OR WHAT ACTS

 **Human** Human, not the author; CI green

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves the colour. The gate itself is a convention today: the ruleset targets no branch (slice 2 gap 3), so the fix sits with the repo admin.

GOVERNANCE TOUCHPOINTS

Code authorship. Authorship is four names: committer, agent co-author, non-requesting approver, accountable owner in the register (module 6.10).

Review load. Diligence lands on the reviewer when the agent owns creation. Count review time beside generation time (modules 6.3, 6.7, 6.9).

MODULES PINNED HERE

6.3 How branches become merges

6.6 Deploy mechanics

6.10 Liability, provenance and code authorship

6.5

Iterate on review comments

 review load

WHAT TO DO

1. Agent answers review comments on the same branch.
2. Reviewer re-reviews until comments are resolved.

WHO OR WHAT ACTS

 **Agent** AI revises, human re-reviews

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Reviewer's comments resolved.

WORKED EXAMPLE

Story B: the same PR regenerates the retailer API docs (green, derived by a build step); publishing them is red and belongs to Stage 8.

COLOUR, AND WHAT WOULD MOVE IT

Amber Stays amber; the realistic shape of the loop. In field data the agent stayed involved in 41 percent of post-merge revisions (slice 1).

GOVERNANCE TOUCHPOINTS

Review load. Diligence lands on the reviewer when the agent owns creation. Count review time beside generation time (modules 6.3, 6.7, 6.9).

MODULES PINNED HERE

6.3 How branches become merges

7 -

Testing and QA

Profile: mixed, with a hard ceiling. 5 tasks.



WHERE THE BASELINE STOOD, AND WHAT THE TESTING FOUND

Baseline amber with the open question 'are the tests correct?'; confirmed, and the question validated by measurement: roughly half of generated tests and assertions are wrong.

SHAPE OF THE STAGE

Agent writes tests, CI runs them, a separate agent attacks them, a human judges whether they are the right tests.

ALSO IN THIS STAGE, HUMAN-LED WITH AI DRAFTING WHERE THE EVIDENCE RULE ALLOWS

Integration and end-to-end tests, accessibility testing (Siteimprove, screen-reader pass), performance and load, contract tests for retailer-facing APIs (story B), exploratory testing, test data with no production data, defect triage, release readiness review.

AS IS AT NBN TODAY

A developer tests their own change, CI runs it, a tester or second developer does a manual pass. Adequacy is judged by coverage. Pain: coverage does not measure correctness, and the reviewer judges tests they did not write against code they did not write. Hypothesis, verify in interviews.

Tasks, in order

7.1

Generate unit tests alongside the change

 token spend

WHAT TO DO

1. Have the agent write tests for every acceptance criterion with the change.
2. Read them. Do not trust green output.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Tests exist for every acceptance criterion.

WORKED EXAMPLE

Slice 2 module 4.

WHO OR WHAT ACTS

 **Agent** AI writes, human reads

COLOUR, AND WHAT WOULD MOVE IT

Amber Stays amber. The green suite that tests nothing is the failure a sub-agent cannot catch because it would write the same test.

GOVERNANCE TOUCHPOINTS

Token spend. Record generating time, reviewing time and tokens consumed for this task (module 6.9). The platform quota is the only hard stop and arrives without warning.

MODULES PINNED HERE

[6.4 How harnesses are used](#)

7.2

Run the CI suite



WHAT TO DO

1. CI runs lint and typecheck, frontend tests, backend tests.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

All checks pass.

WHO OR WHAT ACTS

 **Hook** Hook (CI)

COLOUR, AND WHAT WOULD MOVE IT

Green Green; the four CI jobs are in force in this repo. Green is defensible only because 7.4 exists beside it.

MODULES PINNED HERE

[6.4 How harnesses are used](#)

7.3

Adversarial review pass by a separate agent

 **token spend** **review load**

WHAT TO DO

1. Run the review command with a separate agent, never the one that wrote the code.
2. Surface every finding to a human. Auto-resolve nothing.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Findings surfaced to a human, none auto-resolved.

WORKED EXAMPLE

What the workshop demonstrated with the review command.

WHO OR WHAT ACTS

 **Sub-agent** Sub-agent reviews, human decides on findings

COLOUR, AND WHAT WOULD MOVE IT

Amber Stays amber. An agent may surveil another agent but not itself; it widens coverage and is not an oracle (D1 section 3).

GOVERNANCE TOUCHPOINTS

Token spend. Record generating time, reviewing time and tokens consumed for

this task (module 6.9). The platform quota is the only hard stop and arrives without warning.

Review load. Diligence lands on the reviewer when the agent owns creation. Count review time beside generation time (modules 6.3, 6.7, 6.9).

MODULES PINNED HERE

6.4 How harnesses are used

6.8 Guardrails: what holds and what only shapes

7.4

Judge whether the tests are correct

 review load

WHAT TO DO

1. Inject a fault and see whether the tests catch it (mutation-style).
2. Judge, as a person, whether these are the right tests.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Human confirms the tests would catch a real fault.

WORKED EXAMPLE

CI can say the tests pass; only a person can say they are the right tests.

WHO OR WHAT ACTS

 **Human** Human, with mutation-style tooling as evidence

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. The ceiling is structural: a checking agent has the same next-token properties as the writing agent.

GOVERNANCE TOUCHPOINTS

Review load. Diligence lands on the reviewer when the agent owns creation. Count review time beside generation time (modules 6.3, 6.7, 6.9).

MODULES PINNED HERE

6.4 How harnesses are used

7.5

Security scan and dependency audit



WHAT TO DO

1. CI runs pnpm audit at high severity.
2. A human dispositions any high finding.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

No high-severity finding, or each one dispositioned by a human.

WHO OR WHAT ACTS

Hook Hook (CI)

COLOUR, AND WHAT WOULD MOVE IT

Green Green for the design. In this repo the scan fails red but nothing blocks a merge on it (slice 2 gap 3), so today it is advisory.

MODULES PINNED HERE

[6.4 How harnesses are used](#)

8

Deployment and iteration

Profile: leans autonomous for execution, intervention for judgement. 6 tasks.



WHERE THE BASELINE STOOD, AND WHAT THE TESTING FOUND

Baseline not reached. Classified from slice 1's external evidence, slice 2's deploy mechanics and slice 3's SOCI finding; unvalidated.

SHAPE OF THE STAGE

Preview and roll out behind controllers, a human promotes, AI summarises signals, humans own incidents, and the learning is written back into context for the next cycle.

ALSO IN THIS STAGE, HUMAN-LED WITH AI DRAFTING WHERE THE EVIDENCE RULE ALLOWS

Release notes, change ticket and CAB alignment, retailer notification and API-doc publish (story B), on-call handover, SLO and error-budget review, missed-appointment and rebate reporting (story A), post-release feedback, flag decommissioning.

AS IS AT NBN TODAY

A release is scheduled, approved through a change process, deployed by a pipeline or ops team, monitored by whoever is on call. Post-incident documents are read once. Pain, if confirmed: change process and delivery cadence run on different clocks. Unvalidated; the workshop never reached deployment.

Tasks, in order

8.1

Preview deployment on every PR



WHAT TO DO

1. The hosting platform deploys a preview for every PR and posts the URL.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Preview URL posted on the PR.

WHO OR WHAT ACTS

 **Hook** Hook (CI, hosting platform)

COLOUR, AND WHAT WOULD MOVE IT

Green Green; in force in this repo through the hosting platform.

MODULES PINNED HERE

[6.6 Deploy mechanics](#)

8.2

Progressive rollout



WHAT TO DO

1. Decouple deploy from release with feature flags.
2. Ramp exposure through canary rings.
3. Roll back automatically on a metric threshold, with no human in the loop.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Success-rate and latency thresholds hold at each ring.

WORKED EXAMPLE

Story A: the rebate-relevant change rolls to one ring first and the metric watched is missed-appointment rate, because that is the number the undertaking prices.

WHO OR WHAT ACTS

 **Hook** Hook (rollout controller)

COLOUR, AND WHAT WOULD MOVE IT

Green Green only behind a deterministic controller. This repo has none: merge is deploy and rollback is a person promoting a previous deployment (module 6.6), so today the row is amber in practice.

MODULES PINNED HERE

[6.6 Deploy mechanics](#)

8.3

Production promotion approval



WHAT TO DO

1. Release owner approves promotion, with a rollback plan written first.
2. Where the asset is critical infrastructure, the approval sits inside the written risk management program (SOCI Act).
3. Record the approval against the release.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Approval recorded against the release.

WORKED EXAMPLE

Story B: the retailer notification, the API-doc publish and the deploy are one release event under this approval.

WHO OR WHAT ACTS

 **Human** Human (release owner; regulated where SOCI applies)

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. Google gates it behind role-based approvals; Microsoft requires the rollback plan first.

MODULES PINNED HERE

[6.6 Deploy mechanics](#)

[6.12 Regulatory codification for NBN Co](#)

8.4

Monitoring and alert triage

 review load

WHAT TO DO

1. AI summarises dashboards and correlates alerts.
2. A human dispositions every alert.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Every alert has a human disposition.

WHO OR WHAT ACTS

 **Agent** AI summarises signals, human decides

COLOUR, AND WHAT WOULD MOVE IT

Amber Stays amber. Toward green only for alerts a deterministic rule can already route.

GOVERNANCE TOUCHPOINTS

Review load. Diligence lands on the reviewer when the agent owns creation. Count review time beside generation time (modules 6.3, 6.7, 6.9).

8.5

Incident response and post-incident review



WHAT TO DO

1. Respond.
2. Write the blameless post-incident record.
3. Report incidents where SOCI requires it.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

Blameless post-incident record written.

WHO OR WHAT ACTS

 **Human** Human

COLOUR, AND WHAT WOULD MOVE IT

Red Nothing moves it. Red in every source.

8.6

Fold the learning back into context



WHAT TO DO

1. When a rule, a token or a standard is learned in production, have the agent draft the update to the rules file, skill, design system or ADR.
2. Review and merge it like code.
3. It goes in the repo, not in someone's chat history.

GATE: WHAT HAS TO BE TRUE BEFORE THE NEXT TASK STARTS

The update is merged in the repo.

WORKED EXAMPLE

The brand-colours fix and Team 6's accessibility skill are both instances of this task.

WHO OR WHAT ACTS

 **Agent AI** drafts the update, human approves

COLOUR, AND WHAT WOULD MOVE IT

Amber This task is what backs the migration claim: the write-back is how a stage earns a greener colour next cycle. It closes the loop back to 1.1.

MODULES PINNED HERE

6.7 Collaboration with an agent in the loop

6.11 Prompt practice: written, stored, versioned, shared

— The twelve modules, and where they sit on the map

The module list agreed by the whole team at the 3 Sep reconvene, in D3 order: six build modules from slice 2, six governance modules from slice 3. Each opens into its gate mechanism, the map rows it governs and the one thing it changed in the team's thinking.

6.1

Issue to branch

6.2

How a commit is done

6.3

How branches become merges

6.4

How harnesses are used

6.5

Design-system enforcement in build

6.6

Deploy mechanics

6.7

Collaboration with an agent in the loop

6.8

Guardrails: what holds and what only shapes

6.9

Token accounting and cost

6.10

Liability, provenance and code authorship

6.11

Prompt practice: written, stored, versioned, shared

6.12

Regulatory codification for NBN Co

6.1

Issue to branch

Build (slice 2)

WHAT IT IS

The issue and a committed acceptance-criteria file exist before any agent touches anything; prompts point at the file by name.

GATE MECHANISM (D1 SECTION 3)

Human (Definition of Ready is a checklist a person walks)

MAP ROWS IT GOVERNS

5.3 Write each issue properly

6.1 Agent reads the issue, confirms each acceptance criterion, plans

WHAT IT CHANGED IN THE TEAM'S THINKING

The pattern came from the workshop floor, and it makes Git the shared memory the chat tool does not provide.

6.2

How a commit is done

code authorship

Build (slice 2)

WHAT IT IS

Conventional Commits by a lefthook commit-msg hook; one logical change per commit as a human rule with the reviewer as its check; the co-author trailer as the smallest provenance record.

GATE MECHANISM (D1 SECTION 3)

Hook for format, human for size and boundary

MAP ROWS IT GOVERNS

6.2 Branch, implement, commit

WHAT IT CHANGED IN THE TEAM'S THINKING

Every commit on our main had a human author and no trailer, so the 'live example' was a live example of absence.

6.3

How branches become merges

review load

Build (slice 2)

WHAT IT IS

PR template, four CI jobs, non-author approval, the agent answering review comments on the same branch; what a reviewer checks first on an agent diff.

GATE MECHANISM (D1 SECTION 3)

CI check and human

MAP ROWS IT GOVERNS

6.3 Open a draft PR with description and provenance

6.4 Human code review and merge approval

6.5 Iterate on review comments

WHAT IT CHANGED IN THE TEAM'S THINKING

Branch protection was documented and not in force; the reconciliation starts there.

6.4

How harnesses are used

review load

Build (slice 2)

WHAT IT IS

Two hook layers (git and agent), the deny list, CI, mutation testing as the answer to 'are the tests correct', two sub-agents that are assists and not gates.

GATE MECHANISM (D1 SECTION 3)

Hook, CI check, human for adequacy

MAP ROWS IT GOVERNS

7.1 Generate unit tests alongside the change

7.2 Run the CI suite

7.3 Adversarial review pass by a separate agent

7.4 Judge whether the tests are correct

7.5 Security scan and dependency audit

WHAT IT CHANGED IN THE TEAM'S THINKING

The green suite that tests nothing is the specific failure a sub-agent cannot catch.

6.5

Design-system enforcement in build

Build (slice 2)

WHAT IT IS

Tokens as the rule set, the overlay-diff check as a Playwright visual comparison, the axe scan, the accessibility skill as the amber assist, the regulatory rule ahead of design.

GATE MECHANISM (D1 SECTION 3)

CI check (screenshot diff, axe) and human (the designer reviews a diff, not a screen)

MAP ROWS IT GOVERNS

4.4 Brand and design-system decisions

4.5 Brand compliance check by overlay diff

4.6 Usability and accessibility critique

4.7 Inclusion decisions, including First Nations inclusive-design policy

WHAT IT CHANGED IN THE TEAM'S THINKING

Generation is the wrong tool for enforcing a deterministic system; the fix for the wrong colours went into the context, not the slide.

6.6

Deploy mechanics

Build (slice 2)

WHAT IT IS

Merge is deploy because the hosting platform has no gate of its own; the Firestore rules workflow is out of use; rollback is a person promoting a previous deployment.

GATE MECHANISM (D1 SECTION 3)

Human (merge approval, manual rules deploy, revert) and CI check

MAP ROWS IT GOVERNS

8.1 Preview deployment on every PR

8.2 Progressive rollout

8.3 Production promotion approval

WHAT IT CHANGED IN THE TEAM'S THINKING

The last human decision before production in this repository is the PR approval.

6.7

Collaboration with an agent in the loop

Governance (slice 3)

review load

WHAT IT IS

Four workshop patterns (acceptance criteria as a shared artifact in version control, the screenshot as the specification, skeleton first then pipeline the pages, a review finding encoded as a skill), each with the NBN constraint it has to survive.

GATE MECHANISM (D1 SECTION 3)

Human throughout; the accessibility skill is an assist and the axe scan behind it is the CI check

MAP ROWS IT GOVERNS

2.5 Commit accepted stories and criteria to the repo as Markdown

4.3 Skeleton screens straight from the acceptance criteria

4.6 Usability and accessibility critique

8.6 Fold the learning back into context

WHAT IT CHANGED IN THE TEAM'S THINKING

Parallelism came from the repository, and the team that pipelined its pages paid for it in review time.

6.8

Guardrails: what holds and what only shapes

Governance (slice 3)

WHAT IT IS

Every control states its holder; a document naming a skill or rules file as a control is sent back. Anything that must hold moves into a hook, where it costs no context.

GATE MECHANISM (D1 SECTION 3)

Hook and CI check for pattern-matchable rules, human for judgement, an explicit 'no gate' label for anything that lives only in a prompt

MAP ROWS IT GOVERNS

every holder column

WHAT IT CHANGED IN THE TEAM'S THINKING

The one test a junior needs for any rule in a policy: if the model ignores this, what happens?

6.9

Token accounting and cost

Governance (slice 3)

token spend

WHAT IT IS

Three numbers per task in the spike log: generating time, reviewing time, tokens consumed. Chat during exploration, documents at handover. A register of who owns each AI use and what it consumes.

GATE MECHANISM (D1 SECTION 3)

Human; the platform quota is the only hard stop and it arrives with no warning

MAP ROWS IT GOVERNS

[3.5 Technical spike on the riskiest unknown](#)

WHAT IT CHANGED IN THE TEAM'S THINKING

The certification says nothing about cost, so this is built from the workshop record and the team's own spike; the review-time figure is what makes the decision-load problem visible.

6.10

Liability, provenance and code authorship

Governance (slice 3)

code authorship

WHAT IT IS

Provenance from smallest to largest: the trailer, the confidence card on the PR, the nutrition label that does not exist yet. Authorship as four names. Copyright and licence in agent-written code open for counsel.

GATE MECHANISM (D1 SECTION 3)

Human, backed by one hook the repository could add in an afternoon (commit-msg check refusing an agent commit without a trailer)

MAP ROWS IT GOVERNS

[6.3 Open a draft PR with description and provenance](#)

[6.4 Human code review and merge approval](#)

WHAT IT CHANGED IN THE TEAM'S THINKING

None of the mechanics is switched on in our repository today.

6.11

Prompt practice: written, stored, versioned, shared

Governance (slice 3)

WHAT IT IS

A prompt that matters is a file; the repo has five kinds (CLAUDE.md, skills, agent definitions, the PR template, acceptance files). Files go through the same PR and review as code. Shared as project skill, then plugin.

GATE MECHANISM (D1 SECTION 3)

Human review on the PR and the CI check that already runs on every file

MAP ROWS IT GOVERNS

[5.3 Write each issue properly](#)

[8.6 Fold the learning back into context](#)

WHAT IT CHANGED IN THE TEAM'S THINKING

The thirty-five-file change and the prompt that worked once and was never seen again are the same failure: the absence of the file.

6.12 Regulatory codification for NBN Co

Governance (slice 3)

WHAT IT IS

NBN Co is a corporate Commonwealth entity, so the Commonwealth AI policy (v2.0, 15 Dec 2025) is a model, not an obligation. Seven rules with holders: Privacy Act, register and owner, human oversight, SOCI Act on deploys, DDA and WCAG, First Nations (Stretch RAP VI), Statement of Expectations security.

GATE MECHANISM (D1 SECTION 3)

Human for five of seven rules, hook or CI check for two

MAP ROWS IT GOVERNS

4.7 Inclusion decisions, including First Nations inclusive-design policy

8.3 Production promotion approval

A4

WHAT IT CHANGED IN THE TEAM'S THINKING

The first thing on the list of what binds NBN Co is an absence.

— The stories every example runs on

Story A

technician dispatch. A retailer requests a technician through the maintenance portal; the system matches fault to technician and equipment and schedules the visit. Under the 11 Sep 2018 ACCC enforceable undertaking, a late fix or a missed appointment costs a \$25 rebate each. A scheduling defect is a rebate liability with a regulator behind it.

Story B

retailer notifications. NBN Co is wholesale-only. A service change, scheduled maintenance or an API update has to reach retailers' engineering, sales and legal teams first, with an audit trail. A notification and API-documentation pipeline for retailers.

Base case

the workshop's fault-reporting tool for NBN technicians, which every reader has seen.

Assumed business rules, to verify in the interviews (A4 is an observation with a narrow scope)

RULE	STATEMENT
A1	Every appointment has a committed window; a missed window is a reportable event with a rebate. Assumed, verify in interviews.
A2	Service-affecting changes are notified to affected retailers before the change, with an audit trail. Assumed, verify in interviews.
A3	Some code bases and the retailer-facing API documentation are subject to external review, so their change history has to be readable outside NBN. Assumed, verify in interviews.
A4	At the workshop Alessio asked people not to put the NBN name or company material into the personal AI accounts in use, and not to put company secrets into Claude. That is a rule about the account, not about the tool or the lifecycle. Whether an enterprise account changes it, and which systems if any are off limits to AI, are interview questions. Observed at the workshop, D2 v2 section 9; scope narrowed 8 Sep.

— Where the design and this repository disagree

Green describes the design of the pipeline. Five green or amber rows are not backed by a control in force in [nbn-sdlc-demo](#) today, and each row says so in its "what would move it" line: 4.5 (no Playwright), 6.2 (format step advisory), 6.3 (no trailer hook, no trailers on main), 6.4 (ruleset targets no branch), 7.5 (scan advisory), 8.2 (no rollout controller). These are recorded gaps, not map colours. Stage 8 is unvalidated by the workshop.

D4 for the Planner card "[UX] Finish the lifecycle map". Figma frame B (node 6-155) is the design source for the stage boxes; this page is the operational view and the PDF export in [docs/reports](#) . As-is scenarios are hypotheses from the workshop record, written to be replaced by the developer interviews in Sprint 2.